

Machine Learning Hw4

Ekta Chaudhary

20/04/2020

```
library(ISLR)
library(caret)
library(rpart)
library(rpart.plot)
library(party)
library(partykit)
library(randomForest)
library(ranger)
library(gbm)
library(plotmo)
library(pdp)
library(lime)
library(lasso2)
```

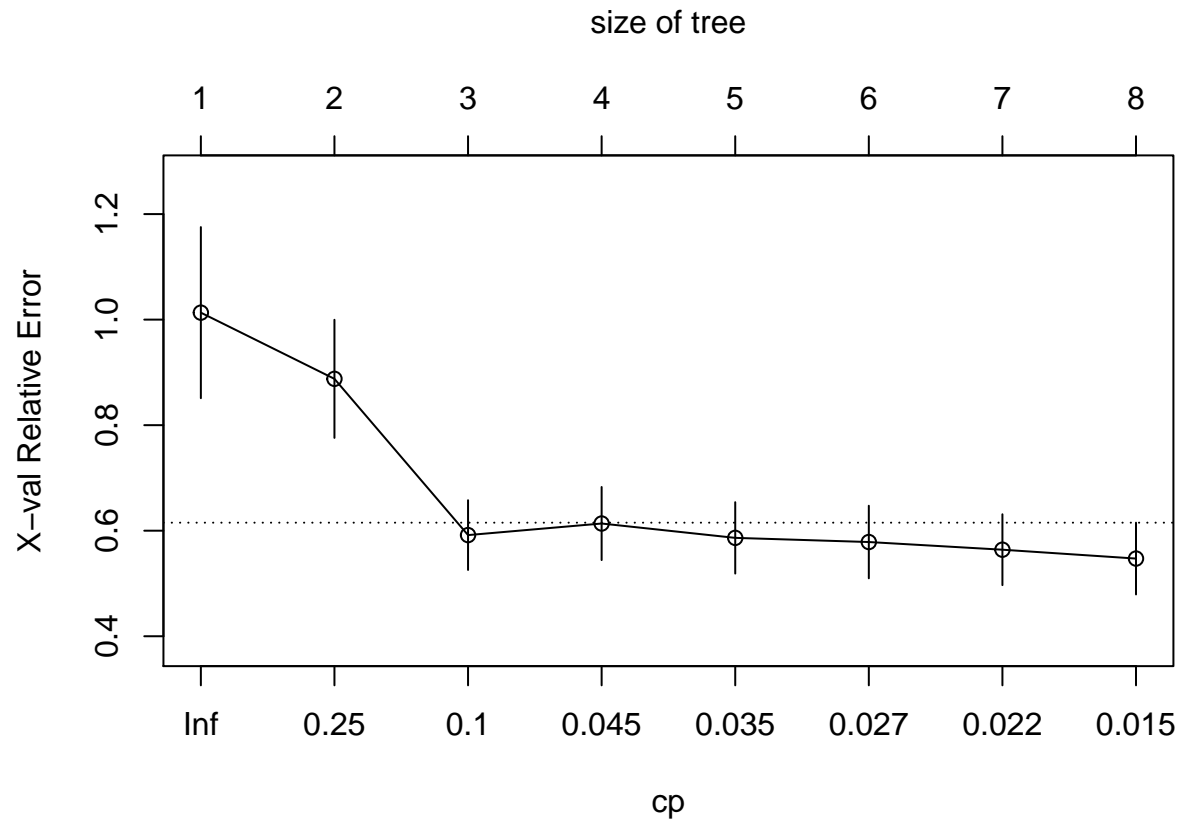
- (a) Fit a regression tree with lpsa as the response and the other variables as predictors. Use cross-validation to determine the optimal tree size. Which tree size corresponds to the lowest cross-validation error? Is this the same as the tree size obtained using the 1 SE rule?

```
set.seed(1)
data("Prostate")
ctrl <- trainControl(method = "cv")
```

```
set.seed(1)
tree <- rpart(formula = lpsa ~ ., data = Prostate,
               control = rpart.control(cp = 0.01))
cpTable <- printcp(tree)
```

```
##
## Regression tree:
## rpart(formula = lpsa ~ ., data = Prostate, control = rpart.control(cp = 0.01))
##
## Variables actually used in tree construction:
## [1] lcavol lweight pgg45
##
## Root node error: 127.92/97 = 1.3187
##
## n= 97
##
##      CP nsplit rel error  xerror   xstd
## 1 0.347108     0  1.00000 1.01323 0.162162
## 2 0.184647     1  0.65289 0.88779 0.111915
## 3 0.059316     2  0.46824 0.59168 0.066102
## 4 0.034756     3  0.40893 0.61359 0.069269
## 5 0.034609     4  0.37417 0.58640 0.067630
## 6 0.021564     5  0.33956 0.57853 0.068772
## 7 0.021470     6  0.31800 0.56398 0.067155
## 8 0.010000     7  0.29653 0.54721 0.068034
```

```
plotcp(tree)
```



```
minErr <- which.min(cpTable[,4])
minErr
```

```
## 8
## 8
```

The tree size 8 corresponds to the lowest cross-validation error.

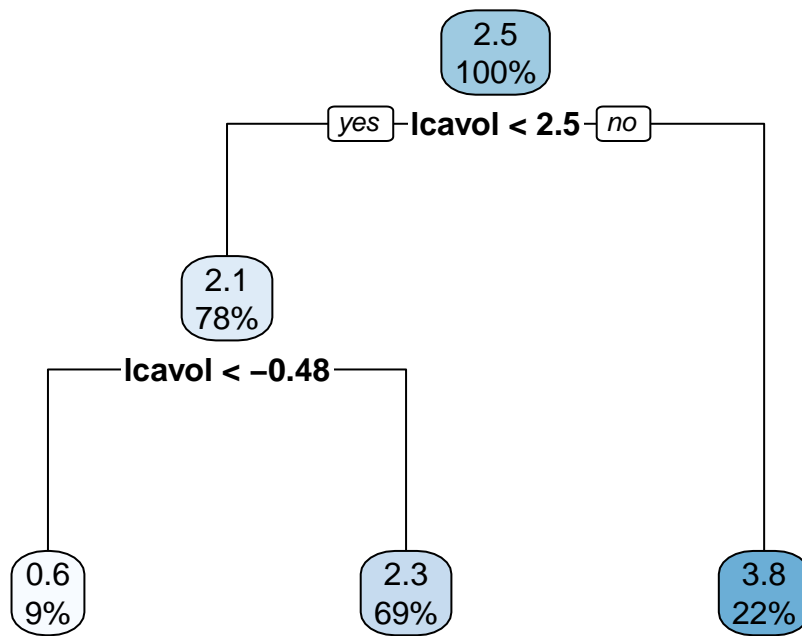
```
cpTable[cpTable[,4] < cpTable[minErr,4] + cpTable[minErr,5],1][1]
```

```
##      3
## 0.05931585
```

The tree size obtained using the 1 SE rule is 3.

- (b) Create a plot of the final tree you choose. Pick one of the terminal nodes, and interpret the information displayed.

```
tree_a = prune(tree, cp = cpTable[cpTable[,4] < cpTable[minErr,4] + cpTable[minErr,5], 1][1])
rpart.plot(tree_a)
```



(c) Perform bagging and report the variable importance

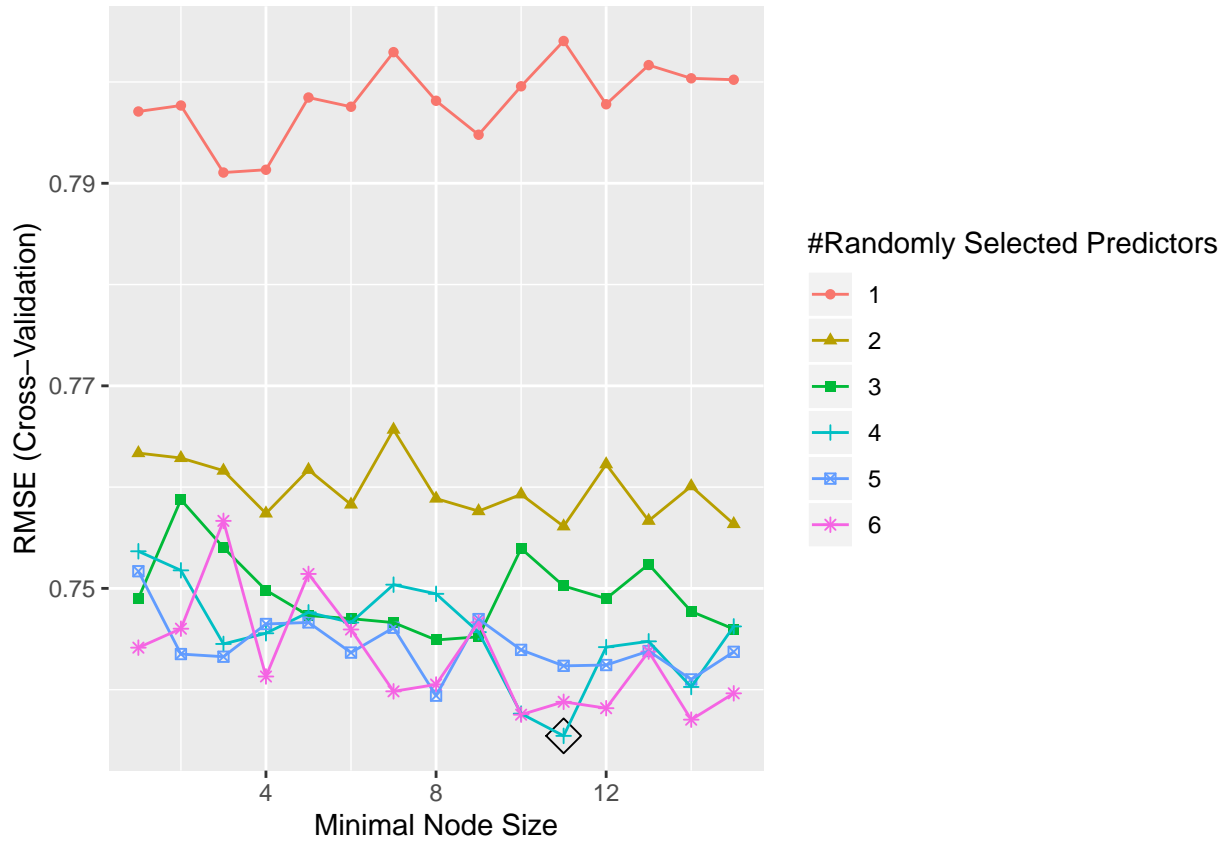
```

bagging.grid <- expand.grid(mtry = 1:6,
                           splitrule = "variance",
                           min.node.size = 1:15)

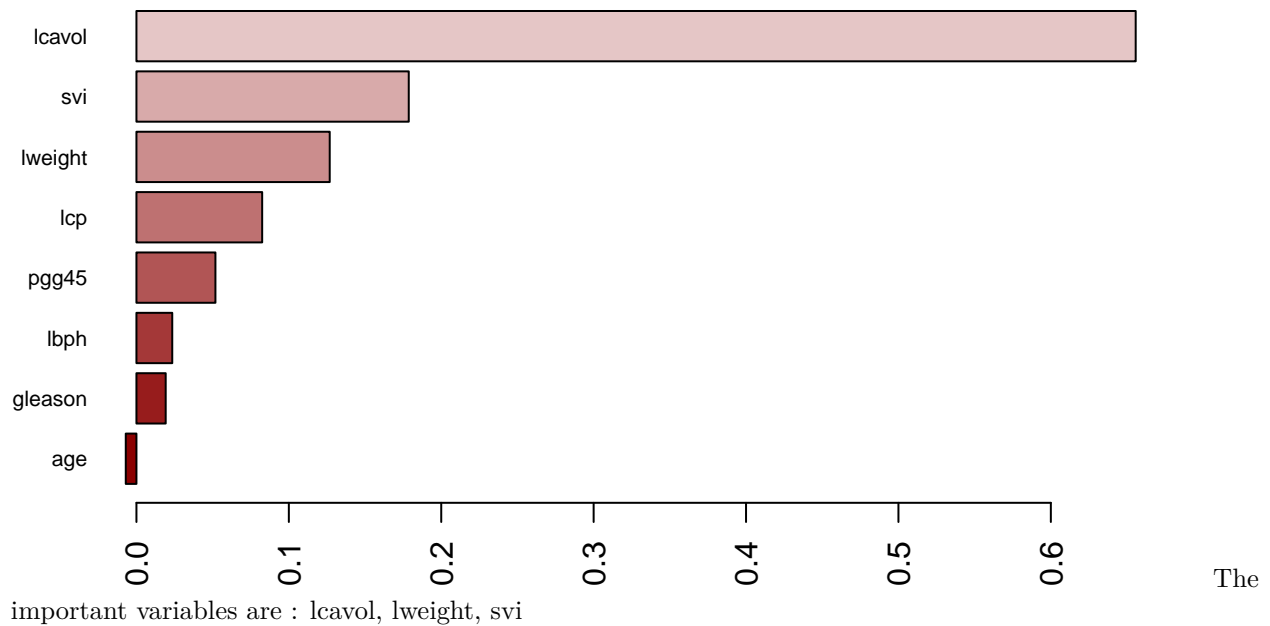
set.seed(1)
bagging <- train(lpsa~., Prostate,
                 method = "ranger",
                 tuneGrid = bagging.grid,
                 trControl = ctrl,
                 importance = "permutation")

ggplot(bagging, highlight = TRUE)

```



```
barplot(sort(ranger::importance(bagging$finalModel), decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(19))
```



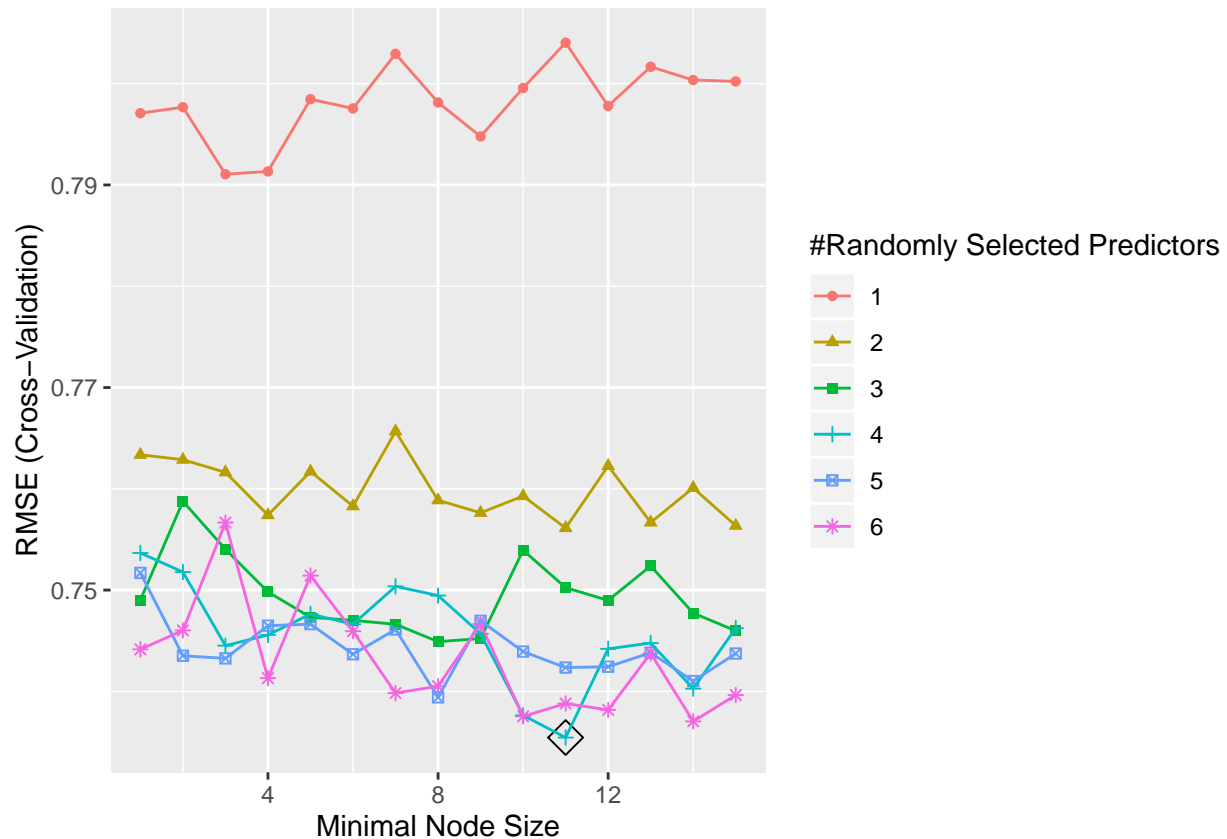
(d) Perform random forests and report the variable importance.

```

rf.grid = expand.grid(mtry = 1:6,
                     splitrule = "variance",
                     min.node.size = 1:15)

set.seed(1)
rf.fit = train(lpsa~., Prostate,
               method = "ranger",
               tuneGrid = rf.grid,
               trControl = ctrl,
               importance = 'permutation')
ggplot(rf.fit, highlight = TRUE)

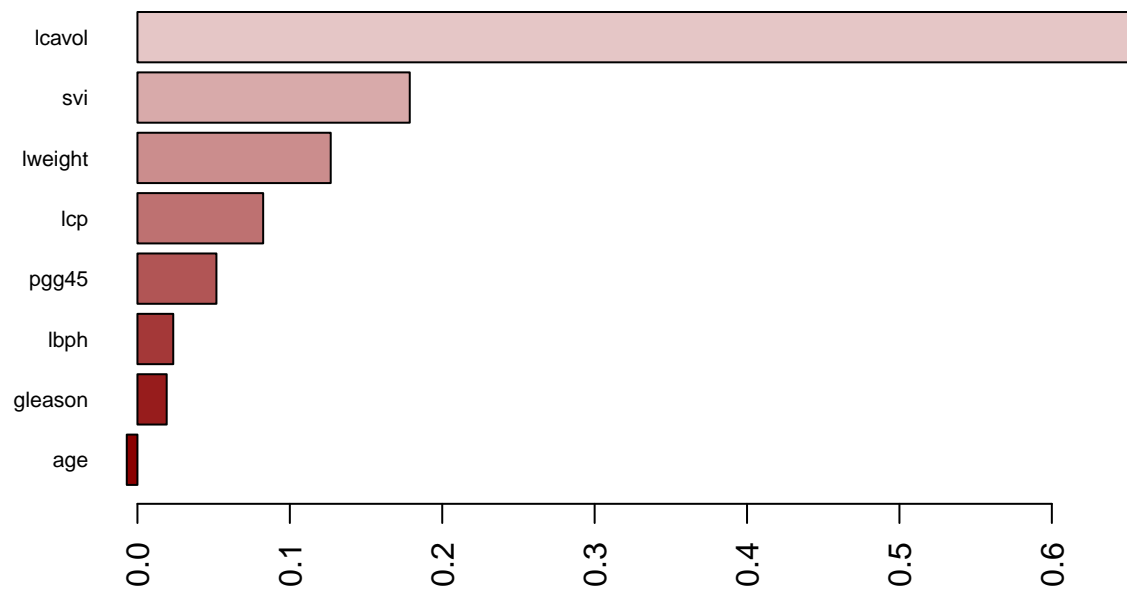
```



```

barplot(sort(ranger::importance(rf.fit$finalModel), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(19))

```

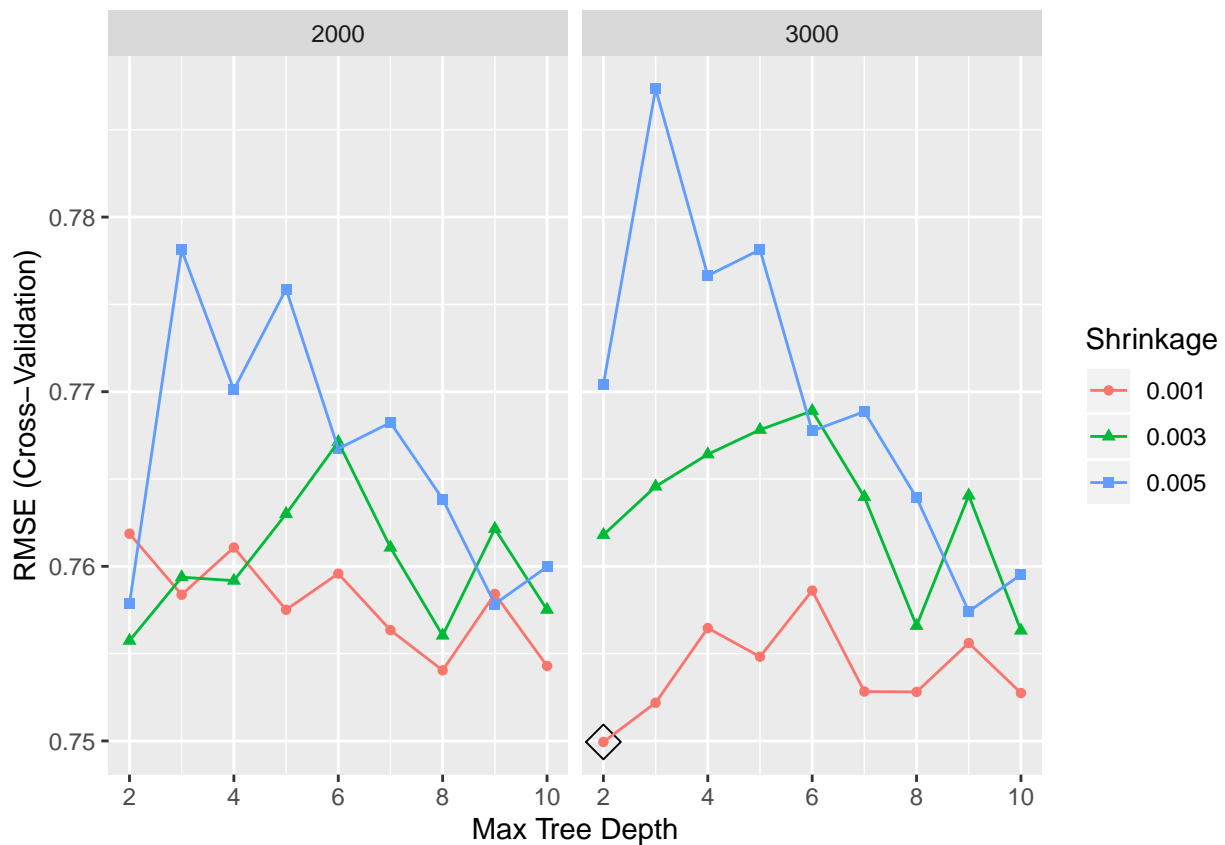


(e) Perform boosting and report the variable importance.

```
gbm.grid <- expand.grid(n.trees = c(2000,3000),
                      interaction.depth = 2:10,
                      shrinkage = c(0.001,0.003,0.005),
                      n.minobsinnode = 1)

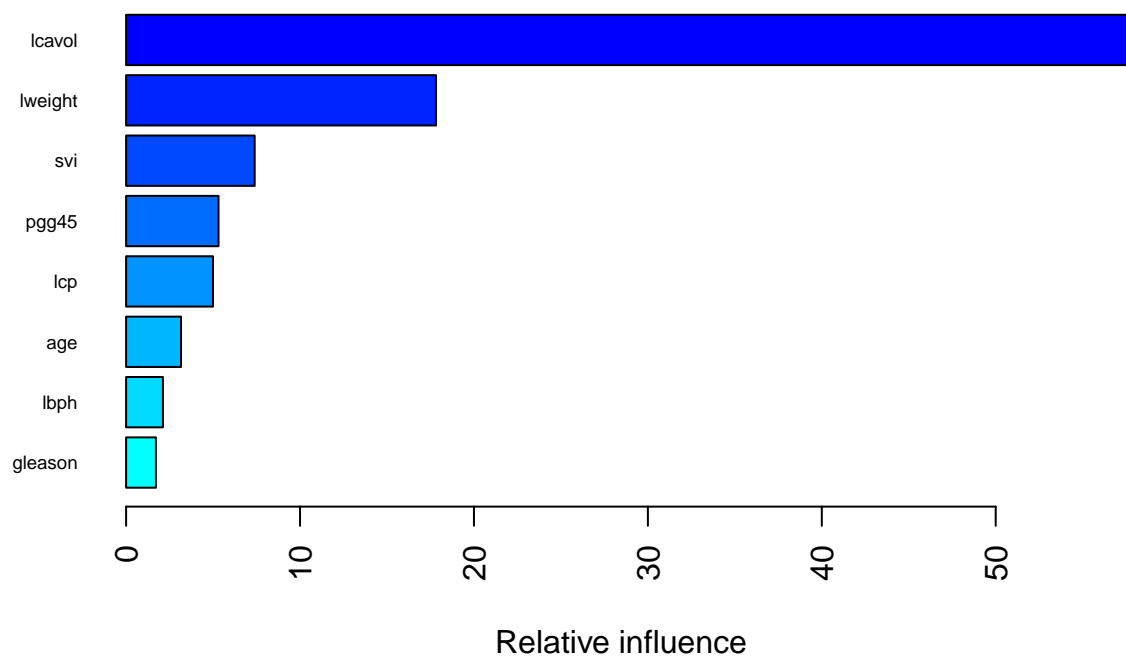
set.seed(1)
gbm.fit <- train(lpsa~., Prostate,
                method = "gbm",
                tuneGrid = gbm.grid,
                trControl = ctrl,
                verbose = FALSE)

ggplot(gbm.fit, highlight = TRUE)
```



Variable importance from boosting can be obtained using the `summary()` function.

```
summary(gbm.fit$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```



```
##      var  rel.inf
```

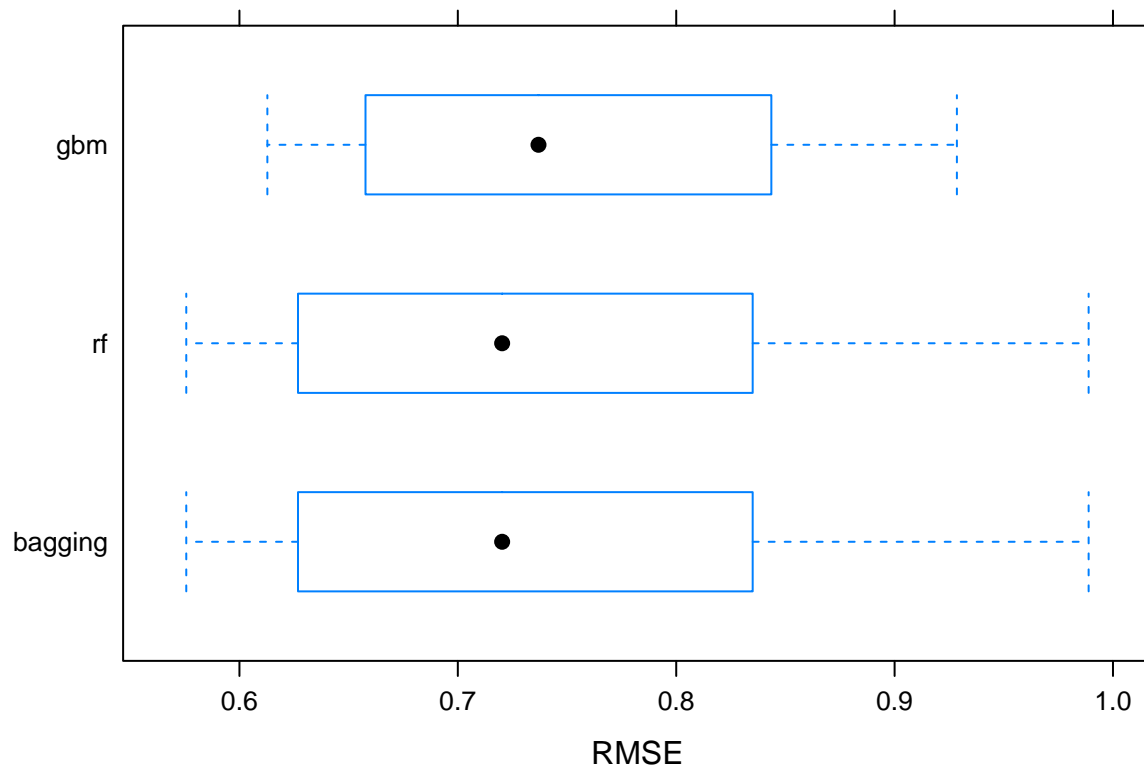
```
## lcavol    lcavol 57.494112
## lweight  lweight 17.818879
## svi      svi    7.392286
## pgg45    pgg45  5.312763
## lcp      lcp    4.994192
## age      age    3.157564
## lbph     lbph   2.115772
## gleason  gleason 1.714432
```

(f) Which of the above models will you select to predict PSA level? Explain.

```
resamp <- resamples(list(bagging = bagging, rf = rf.fit, gbm = gbm.fit))
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: bagging, rf, gbm
## Number of resamples: 10
##
## MAE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## bagging 0.4794948 0.5357349 0.6115087 0.5980269 0.6632726 0.7001466    0
## rf      0.4794948 0.5357349 0.6115087 0.5980269 0.6632726 0.7001466    0
## gbm     0.4988605 0.5382599 0.6238535 0.6060614 0.6541746 0.7189084    0
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## bagging 0.5756510 0.6319992 0.7202955 0.7354396 0.8235912 0.9888492    0
## rf      0.5756510 0.6319992 0.7202955 0.7354396 0.8235912 0.9888492    0
## gbm     0.6127814 0.6613586 0.7369285 0.7499428 0.8330904 0.9285269    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## bagging 0.3250728 0.5203503 0.6199793 0.6083602 0.7499733 0.7763990    0
## rf      0.3250728 0.5203503 0.6199793 0.6083602 0.7499733 0.7763990    0
## gbm     0.3912307 0.4940316 0.6253058 0.6103438 0.7211938 0.8171347    0
```

```
bwplot(resamp, metric = "RMSE")
```

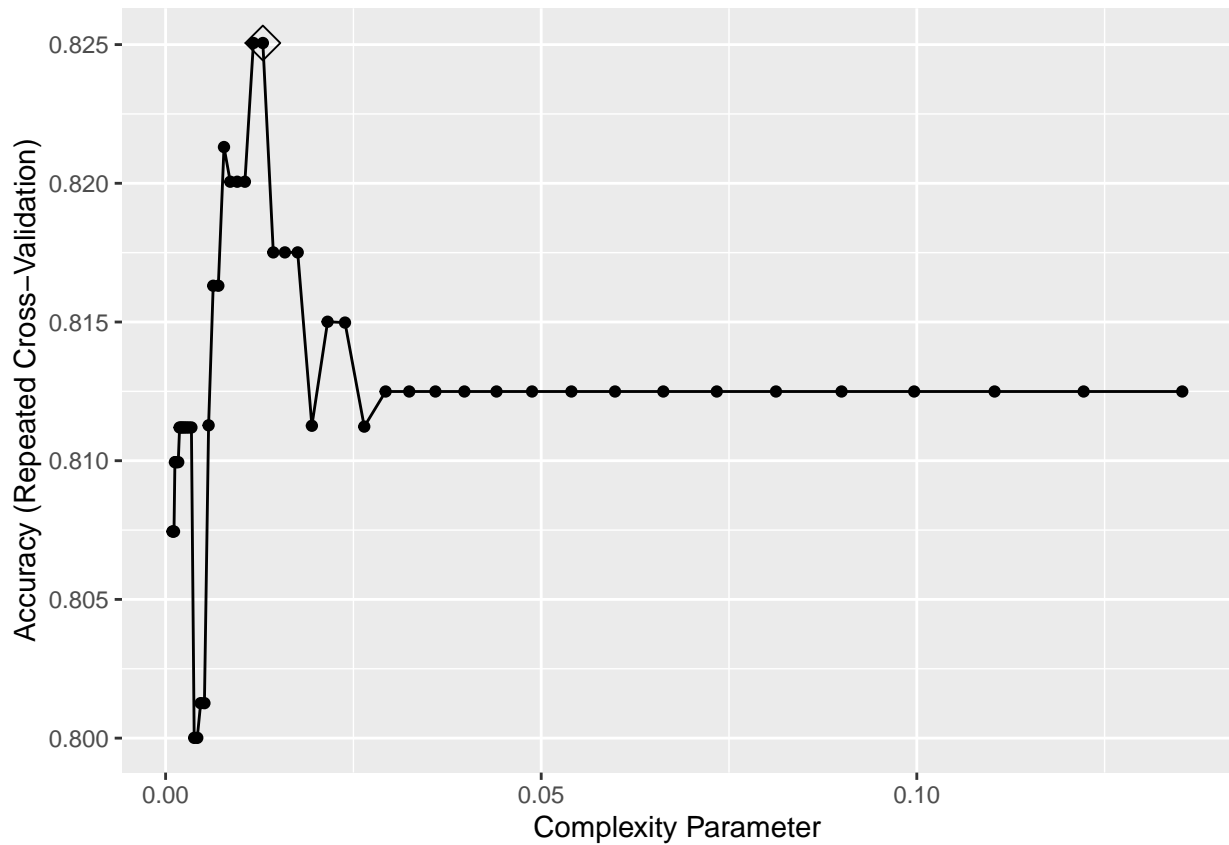



2. This problem involves the OJ data in the ISLR package. The data contains 1070 purchases where the customers either purchased Citrus Hill or Minute Maid Orange Juice. A number of characteristics of customers and products are recorded. Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations. Use `set.seed()` for reproducible results.

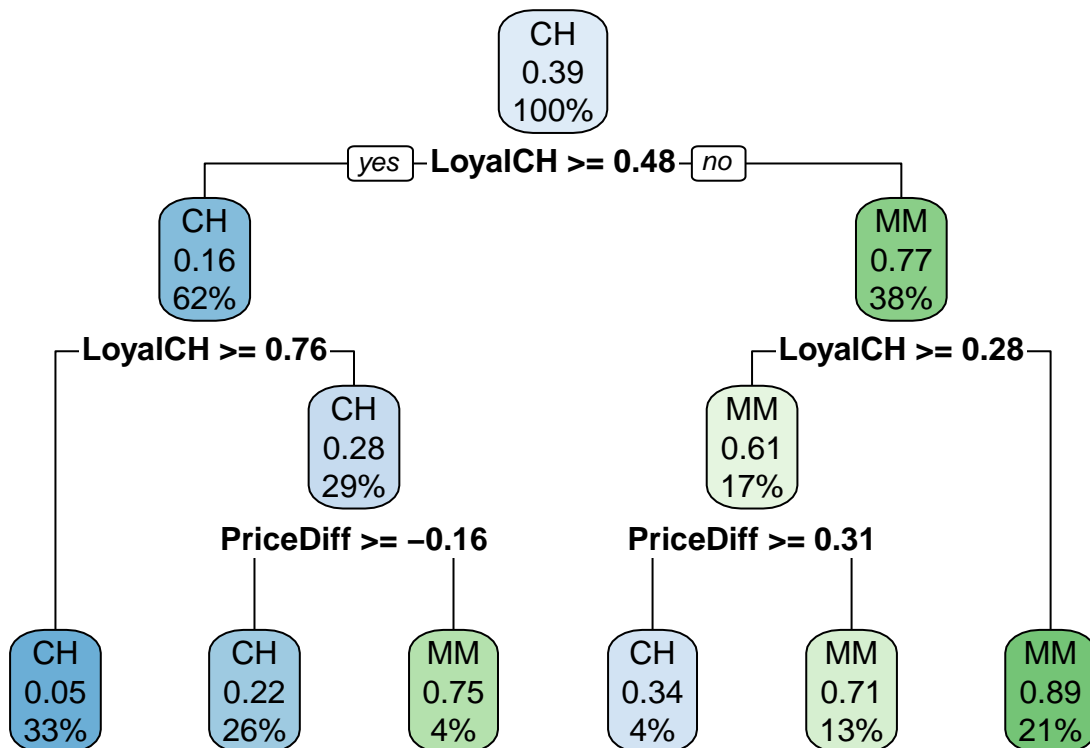
```
data("OJ")
set.seed(1)
rowTrain = createDataPartition(y = OJ$Purchase,
                               p = 0.747,
                               list = FALSE)
ctrl <- trainControl(method = "repeatedcv")
```

- (a) Fit a classification tree to the training set, with Purchase as the response and the other variables as predictors. Use cross-validation to determine the tree size and create a plot of the final tree. Predict the response on the test data. What is the test classification error rate?

```
set.seed(1)
rpart.class <- train(Purchase ~., OJ,
                    subset = rowTrain,
                    method = "rpart",
                    tuneGrid = data.frame(cp = exp(seq(-7, -2, len = 50))),
                    trControl = ctrl,
                    metric = "Accuracy")
ggplot(rpart.class, highlight = T)
```



```
rpart.plot(rpart.class$finalModel)
```



```
rpart.class$bestTune
```

```
##           cp  
## 27 0.01294638
```

```
rpart.pred <- predict(rpart.class, newdata = OJ[-rowTrain,])  
confusionMatrix(rpart.pred,  
                 reference = OJ$Purchase[-rowTrain])
```

```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction  CH  MM  
##           CH 142  24  
##           MM  23  81  
##  
##           Accuracy : 0.8259  
##           95% CI : (0.7753, 0.8692)  
## No Information Rate : 0.6111  
## P-Value [Acc > NIR] : 1.626e-14  
##  
##           Kappa : 0.6331  
##  
## McNemar's Test P-Value : 1  
##  
##           Sensitivity : 0.8606  
##           Specificity : 0.7714  
##           Pos Pred Value : 0.8554  
##           Neg Pred Value : 0.7788  
##           Prevalence : 0.6111  
##           Detection Rate : 0.5259  
## Detection Prevalence : 0.6148  
##           Balanced Accuracy : 0.8160  
##  
##           'Positive' Class : CH  
##
```

```
# Error rate  
error_rate = mean(rpart.pred != OJ$Purchase[-rowTrain]) * 100  
cat(c("The error rate for the classification tree is", error_rate, '%'))
```

```
## The error rate for the classification tree is 17.4074074074074 %
```

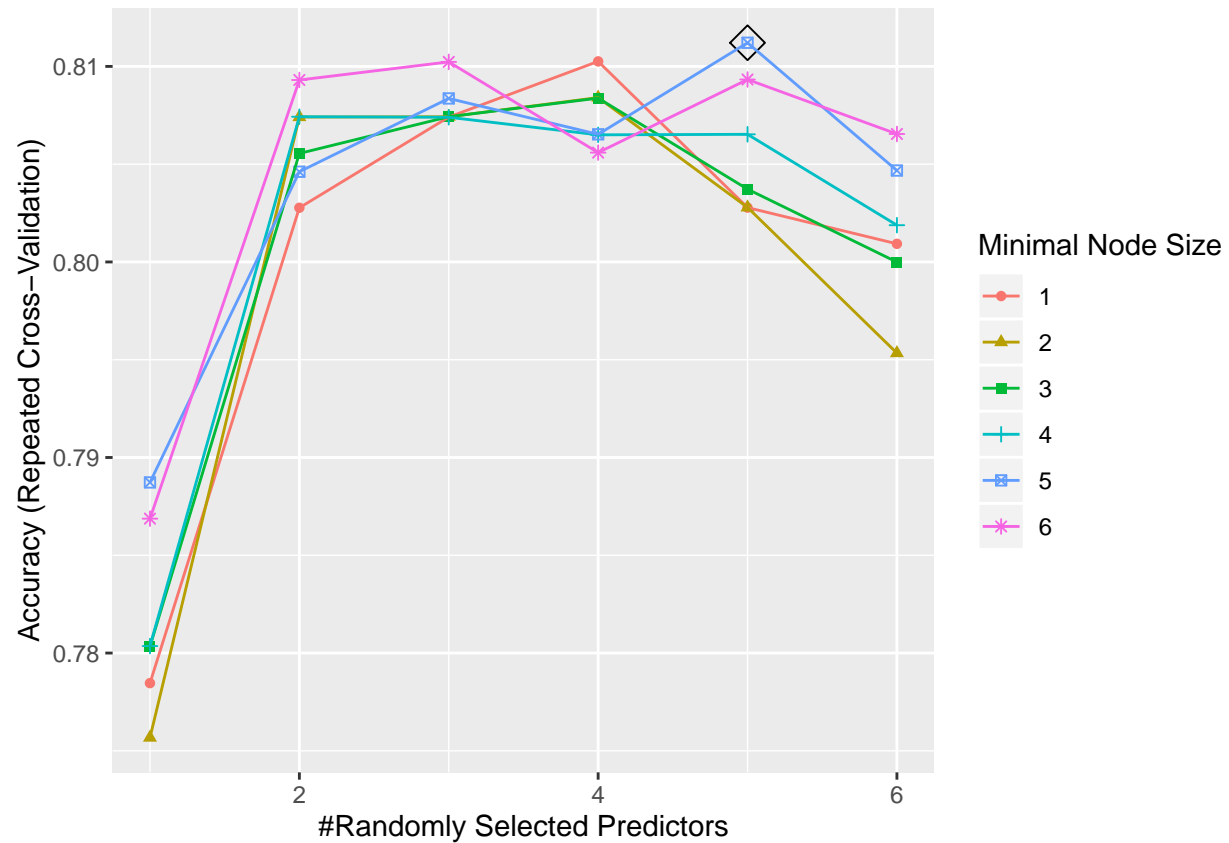
- (b) Perform random forests on the training set and report variable importance. What is the test error rate?

```
rf.grid1 = expand.grid(mtry = 1:6,  
                      splitrule = "gini",  
                      min.node.size = 1:6)  
set.seed(1)  
rf.fit1 = train(Purchase~., OJ,
```

```

method = "ranger",
tuneGrid = rf.grid1,
trControl = ctrl,
importance = 'permutation')
ggplot(rf.fit1, highlight = TRUE)


```



```

barplot(sort(ranger::importance(rf.fit1$finalModel), decreasing = FALSE),
las = 2, horiz = TRUE, cex.names = 0.7,
col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(19))


```



(c) Perform boosting on the training set and report variable importance. What is the test error rate?


```
gbm2.grid <- expand.grid(n.trees = c(2000,3000,4000),
                        interaction.depth = 1:6,
                        shrinkage = c(0.001,0.003,0.005),
                        n.minobsinnode = 1)

set.seed(1)
gbm2.fit <- train(Purchase ~., OJ,
                  subset = rowTrain,
                  tuneGrid = gbm2.grid,
                  trControl = ctrl,
                  method = "gbm",
                  distribution = "adaboost",
                  metric = "Accuracy",
                  verbose = FALSE)
ggplot(gbm2.fit, highlight = TRUE)
```



The selected model had a maximum tree depth of 3 and learning rate of 0.003 with 3000 boosting iterations.

```
summary(gbm2.fit$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```



```
##          var    rel.inf
```

## LoyalCH	LoyalCH	64.2497567
## PriceDiff	PriceDiff	10.7738464
## WeekofPurchase	WeekofPurchase	6.3868678
## ListPriceDiff	ListPriceDiff	3.5154509
## SalePriceMM	SalePriceMM	2.9987268
## StoreID	StoreID	2.4774727
## STORE	STORE	2.2892477
## PriceMM	PriceMM	1.4530011
## SpecialCH	SpecialCH	1.1007800
## DiscMM	DiscMM	0.9273452
## SalePriceCH	SalePriceCH	0.8540572
## DiscCH	DiscCH	0.7788827
## PriceCH	PriceCH	0.7777353
## PctDiscMM	PctDiscMM	0.5350004
## SpecialMM	SpecialMM	0.4903835
## PctDiscCH	PctDiscCH	0.2340908
## Store7Yes	Store7Yes	0.1573546